

Multiple Linear Regression is one of the important regression algorithms which models the linear relationship between a single dependent continuous variable and more than one independent variable.

```
In [ ]: """Problem Description:

We have a dataset of 50 start-up companies.
This dataset contains five main information:
R&D Spend, Administration Spend, Marketing Spend, State, and Profit for a financial year.
Our goal is to create a model that can easily determine which company has a maximum profit,
and which is the most affecting factor for the profit of a company."""
```

```
Out[1]: 'Problem Description:\n\nWe have a dataset of 50 start-up companies. \nThis dataset contains five main information: \nR&D Spend, Administration Spend, Marketing Spend, State, and Profit for a financial year. \nOur goal is to create a model that can easily determine which company has a maximum profit, and which is the most affecting factor for the profit of a company.'
```

```
In [ ]: #Steps to implement Machine Learning
"""duplicated Values Remove
Null values Remove
Categorical Column remove
standardised data
splitting the data
train the data
testing the data"""
```

```
Out[2]: 'duplicated Values Remove\nNull values Remove\nCategorical Column remove\nstandardised data\nsplitting the data\ntrain the data\ntesting the data'
```

```
In [ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [ ]: df=pd.read_csv("/content/sample_data/35 Startups_Multiple_Linear_Regression")
```

```
In [ ]: df.head()
```

```
Out[6]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
In [ ]: df.shape
```

```
Out[8]: (50, 5)
```

```
In [ ]: df.size
```

```
Out[9]: 250
```

```
In [ ]: df.dtypes
```

```
Out[10]: R&D Spend      float64
Administration  float64
Marketing Spend  float64
State            object
Profit           float64
dtype: object
```

```
In [ ]: df.describe()
```

```
Out[11]:
```

	R&D Spend	Administration	Marketing Spend	Profit
count	50.000000	50.000000	50.000000	50.000000
mean	73721.615600	121344.639600	211025.097800	112012.639200
std	45902.256482	28017.802755	122290.310726	40306.180338
min	0.000000	51283.140000	0.000000	14681.400000
25%	39936.370000	103730.875000	129300.132500	90138.902500
50%	73051.080000	122699.795000	212716.240000	107978.190000
75%	101602.800000	144842.180000	299469.085000	139765.977500
max	165349.200000	182645.560000	471784.100000	192261.830000

```
In [ ]: df.columns
```

```
Out[12]: Index(['R&D Spend', 'Administration', 'Marketing Spend', 'State', 'Profit'], dtype='object')
```

```
In [ ]: df.nunique()
```

```
Out[13]: R&D Spend      49
Administration  50
Marketing Spend  48
State            3
Profit           50
dtype: int64
```

```
In [ ]: df.duplicated()
```

```
Out[14]: 0      False
1      False
2      False
3      False
4      False
5      False
6      False
7      False
8      False
9      False
10     False
11     False
12     False
13     False
14     False
15     False
16     False
17     False
18     False
19     False
20     False
21     False
22     False
23     False
24     False
25     False
26     False
27     False
28     False
29     False
30     False
31     False
32     False
33     False
34     False
35     False
36     False
37     False
38     False
39     False
40     False
41     False
42     False
43     False
44     False
45     False
46     False
47     False
48     False
49     False
dtype: bool
```

```
In [ ]: df.duplicated().sum()
```

```
Out[15]: 0
```

```
In [ ]: df.drop_duplicates(inplace=True)
```

```
In [ ]: df.duplicated().sum()
```

```
Out[17]: 0
```

```
In [ ]: df.dropna() #dropping NA values
```

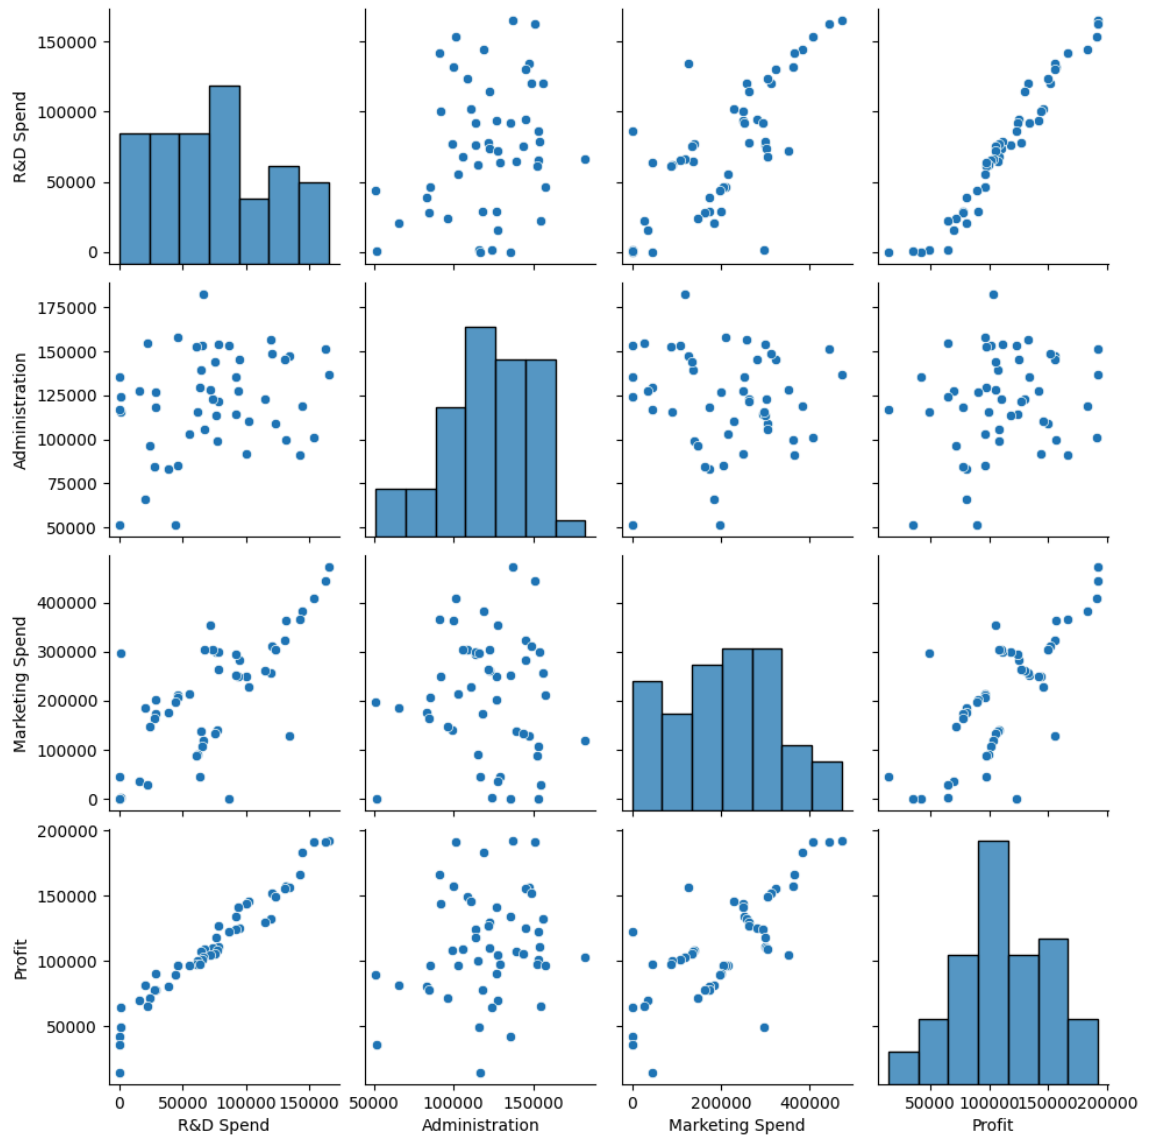
Out[18]:

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84
33	55493.95	103057.49	214634.81	Florida	96778.92
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06

	R&D Spend	Administration	Marketing Spend	State	Profit
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41
49	0.00	116983.80	45173.06	California	14681.40

```
In [ ]: sns.pairplot(df)
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x7b6f8ae9a2f0>
```



```
In [ ]: X=df[["R&D Spend","Administration","Marketing Spend"]]
```

```
In [ ]: X.shape
```

```
Out[23]: (50, 3)
```

```
In [ ]: from sklearn.preprocessing import StandardScaler #scaleing the features bec
```

```
In [ ]: ss=StandardScaler()  
X=ss.fit_transform(X)
```



```
In [ ]: X
```

```
Out[32]: array([[ 2.01641149e+00,  5.60752915e-01,  2.15394309e+00],
 [ 1.95586034e+00,  1.08280658e+00,  1.92360040e+00],
 [ 1.75436374e+00, -7.28257028e-01,  1.62652767e+00],
 [ 1.55478369e+00, -9.63646307e-02,  1.42221024e+00],
 [ 1.50493720e+00, -1.07991935e+00,  1.28152771e+00],
 [ 1.27980001e+00, -7.76239071e-01,  1.25421046e+00],
 [ 1.34006641e+00,  9.32147208e-01, -6.88149930e-01],
 [ 1.24505666e+00,  8.71980011e-01,  9.32185978e-01],
 [ 1.03036886e+00,  9.86952101e-01,  8.30886909e-01],
 [ 1.09181921e+00, -4.56640246e-01,  7.76107440e-01],
 [ 6.20398248e-01, -3.87599089e-01,  1.49807267e-01],
 [ 5.93085418e-01, -1.06553960e+00,  3.19833623e-01],
 [ 4.43259872e-01,  2.15449064e-01,  3.20617441e-01],
 [ 4.02077603e-01,  5.10178953e-01,  3.43956788e-01],
 [ 1.01718075e+00,  1.26919939e+00,  3.75742273e-01],
 [ 8.97913123e-01,  4.58678535e-02,  4.19218702e-01],
 [ 9.44411957e-02,  9.11841968e-03,  4.40446224e-01],
 [ 4.60720127e-01,  8.55666318e-01,  5.91016724e-01],
 [ 3.96724938e-01, -2.58465367e-01,  6.92992062e-01],
 [ 2.79441650e-01,  1.15983657e+00, -1.74312698e+00],
 [ 5.57260867e-02, -2.69587651e-01,  7.23925995e-01],
 [ 1.02723599e-01,  1.16918609e+00,  7.32787791e-01],
 [ 6.00657792e-03,  5.18495648e-02,  7.62375876e-01],
 [-1.36200724e-01, -5.62211268e-01,  7.74348908e-01],
 [ 7.31146008e-02, -7.95469167e-01, -5.81939297e-01],
 [-1.99311688e-01,  6.56489139e-01, -6.03516725e-01],
 [ 3.53702028e-02,  8.21717916e-01, -6.35835495e-01],
 [-3.55189938e-02,  2.35068543e-01,  1.17427116e+00],
 [-1.68792717e-01,  2.21014050e+00, -7.67189437e-01],
 [-1.78608540e-01,  1.14245677e+00, -8.58133663e-01],
 [-2.58074369e-01, -2.05628659e-01, -9.90357166e-01],
 [-2.76958231e-01,  1.13055391e+00, -1.01441945e+00],
 [-2.26948675e-01,  2.83923813e-01, -1.36244978e+00],
 [-4.01128925e-01, -6.59324033e-01,  2.98172434e-02],
 [-6.00682122e-01,  1.31053525e+00, -1.87861793e-03],
 [-6.09749941e-01, -1.30865753e+00, -4.54931587e-02],
 [-9.91570153e-01,  2.05924691e-01, -8.17625734e-02],
 [-6.52532310e-01, -2.52599402e+00, -1.15608256e-01],
 [-1.17717755e+00, -1.99727037e+00, -2.12784866e-01],
 [-7.73820359e-01, -1.38312156e+00, -2.97583276e-01],
 [-9.89577015e-01, -1.00900218e-01, -3.15785883e-01],
 [-1.00853372e+00, -1.32079581e+00, -3.84552407e-01],
 [-1.10210556e+00, -9.06937535e-01, -5.20595959e-01],
 [-1.28113364e+00,  2.17681524e-01, -1.44960468e+00],
 [-1.13430539e+00,  1.20641936e+00, -1.50907418e+00],
 [-1.60035036e+00,  1.01253936e-01, -1.72739998e+00],
 [-1.59341322e+00, -1.99321741e-01,  7.11122474e-01],
 [-1.62236202e+00,  5.07721876e-01, -1.74312698e+00],
 [-1.61043334e+00, -2.50940884e+00, -1.74312698e+00],
 [-1.62236202e+00, -1.57225506e-01, -1.36998473e+00]])
```

```
In [ ]: Y=df["Profit"]
```

```
In [ ]: Y.shape
```

```
Out[35]: (50,)
```

```
In [ ]: from sklearn.model_selection import train_test_split
```

```
In [ ]: #split the dataset into training data and training data
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.25,random_st
```

```
In [ ]: X_train.shape,X_test.shape,Y_train.shape,Y_test.shape
```

```
Out[38]: ((37, 3), (13, 3), (37,), (13,))
```

```
In [ ]: #now implement the Linear Regression Model
```

```
In [ ]: from sklearn.linear_model import LinearRegression
```

```
In [ ]: lr=LinearRegression()
```

```
In [ ]: lr.fit(X_train,Y_train)
```

```
Out[42]: LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [ ]: y_pred=lr.predict(X_test)
```

```
In [ ]: y_pred
```

```
Out[44]: array([ 88361.6924659 , 109068.75037541,  66233.18132181,  70645.38100143,
 48118.47333074, 115786.66944536, 171799.96557761,  99617.55808099,
159031.78297409, 157877.26074356,  83222.30531514, 179714.94106163,
 75105.99525989])
```

```
In [ ]: lr.score(X_train,Y_train)
```

```
Out[45]: 0.9310605936487033
```

```
In [ ]: lr.score(X_test,Y_test)
```

```
Out[46]: 0.9878392927652377
```

```
In [ ]: lr.intercept_
```

```
Out[47]: 111675.78913907126
```

```
In [ ]: lr.coef_
```

```
Out[48]: array([36077.48534245, -219.17000538,  2819.81544887])
```

```
In [ ]: #now find the error/ cost function/loss function
```

```
In [ ]: from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_e
```

```
In [ ]: mean_absolute_error(Y_test,y_pred)
```

Out[51]: 3476.6285513802627

```
In [ ]: mean_absolute_percentage_error(Y_test,y_pred)
```

Out[52]: 0.032672215269165125

Method 2 from sir point of views- refer from project no 8

```
In [ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [ ]: df2=pd.read_csv("/content/sample_data/35 Startups_Multiple_Linear_Regression
```

```
In [ ]: df2.head(5)
```

Out[4]:

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
In [ ]: df2.shape
```

Out[5]: (50, 5)

```
In [ ]: df2.size
```

Out[6]: 250

```
In [ ]: df2.dtypes
```

Out[7]:

R&D Spend	float64
Administration	float64
Marketing Spend	float64
State	object
Profit	float64
dtype:	object

```
In [ ]: df2.ndim
```

Out[8]: 2

```
In [ ]: df2.describe()
```

```
Out[10]:
```

	R&D Spend	Administration	Marketing Spend	Profit
count	50.000000	50.000000	50.000000	50.000000
mean	73721.615600	121344.639600	211025.097800	112012.639200
std	45902.256482	28017.802755	122290.310726	40306.180338
min	0.000000	51283.140000	0.000000	14681.400000
25%	39936.370000	103730.875000	129300.132500	90138.902500
50%	73051.080000	122699.795000	212716.240000	107978.190000
75%	101602.800000	144842.180000	299469.085000	139765.977500
max	165349.200000	182645.560000	471784.100000	192261.830000

```
In [ ]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   R&D Spend              50 non-null    float64
1   Administration         50 non-null    float64
2   Marketing Spend        50 non-null    float64
3   State                  50 non-null    object
4   Profit                  50 non-null    float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

```
In [ ]: df2.columns
```

```
Out[16]: Index(['R&D Spend', 'Administration', 'Marketing Spend', 'State', 'Profit'], dtype='object')
```

```
In [ ]: df2.nunique()
```

```
Out[17]: R&D Spend          49
Administration          50
Marketing Spend          48
State                    3
Profit                   50
dtype: int64
```

```
In [ ]: df2["R&D Spend"].unique()
```

```
Out[18]: array([165349.2 , 162597.7 , 153441.51, 144372.41, 142107.34, 131876.9 ,
        134615.46, 130298.13, 120542.52, 123334.88, 101913.08, 100671.96,
        93863.75,  91992.39, 119943.24, 114523.61,  78013.11,  94657.16,
        91749.16,  86419.7 ,  76253.86,  78389.47,  73994.56,  67532.53,
        77044.01,  64664.71,  75328.87,  72107.6 ,  66051.52,  65605.48,
        61994.48,  61136.38,  63408.86,  55493.95,  46426.07,  46014.02,
        28663.76,  44069.95,  20229.59,  38558.51,  28754.33,  27892.92,
        23640.93,  15505.73,  22177.74,  1000.23,  1315.46,    0. ,
        542.05])
```

```
In [ ]: df2["Administration"].unique()
```

```
Out[19]: array([136897.8 , 151377.59, 101145.55, 118671.85,  91391.77,  99814.71,
        147198.87, 145530.06, 148718.95, 108679.17, 110594.11,  91790.61,
        127320.38, 135495.07, 156547.42, 122616.84, 121597.55, 145077.58,
        114175.79, 153514.11, 113867.3 , 153773.43, 122782.75, 105751.03,
         99281.34, 139553.16, 144135.98, 127864.55, 182645.56, 153032.06,
        115641.28, 152701.92, 129219.61, 103057.49, 157693.92,  85047.44,
        127056.21,  51283.14,  65947.93,  82982.09, 118546.05,  84710.77,
         96189.63, 127382.3 , 154806.14, 124153.04, 115816.21, 135426.92,
        51743.15, 116983.8 ])
```

```
In [ ]: df2["Marketing Spend"].unique()
```

```
Out[20]: array([471784.1 , 443898.53, 407934.54, 383199.62, 366168.42, 362861.36,
        127716.82, 323876.68, 311613.29, 304981.62, 229160.95, 249744.55,
        249839.44, 252664.93, 256512.92, 261776.23, 264346.06, 282574.31,
        294919.57,    0. , 298664.47, 299737.29, 303319.26, 304768.73,
        140574.81, 137962.62, 134050.07, 353183.81, 118148.2 , 107138.38,
         91131.24,  88218.23,  46085.25, 214634.81, 210797.67, 205517.64,
        201126.82, 197029.42, 185265.1 , 174999.3 , 172795.67, 164470.71,
        148001.11,  35534.17,  28334.72,   1903.93, 297114.46,  45173.06])
```

```
In [ ]: df2["State"].unique()
```

```
Out[21]: array(['New York', 'California', 'Florida'], dtype=object)
```

```
In [ ]: df2["Profit"].unique()
```

```
Out[22]: array([192261.83, 191792.06, 191050.39, 182901.99, 166187.94, 156991.12,
        156122.51, 155752.6 , 152211.77, 149759.96, 146121.95, 144259.4 ,
        141585.52, 134307.35, 132602.65, 129917.04, 126992.93, 125370.37,
        124266.9 , 122776.86, 118474.03, 111313.02, 110352.25, 108733.99,
        108552.04, 107404.34, 105733.54, 105008.31, 103282.38, 101004.64,
         99937.59,  97483.56,  97427.84,  96778.92,  96712.8 ,  96479.51,
         90708.19,  89949.14,  81229.06,  81005.76,  78239.91,  77798.83,
         71498.49,  69758.98,  65200.33,  64926.08,  49490.75,  42559.73,
        35673.41,  14681.4 ])
```

```
In [ ]: df2.duplicated()
```

```
Out[25]: 0      False
         1      False
         2      False
         3      False
         4      False
         5      False
         6      False
         7      False
         8      False
         9      False
        10      False
        11      False
        12      False
        13      False
        14      False
        15      False
        16      False
        17      False
        18      False
        19      False
        20      False
        21      False
        22      False
        23      False
        24      False
        25      False
        26      False
        27      False
        28      False
        29      False
        30      False
        31      False
        32      False
        33      False
        34      False
        35      False
        36      False
        37      False
        38      False
        39      False
        40      False
        41      False
        42      False
        43      False
        44      False
        45      False
        46      False
        47      False
        48      False
        49      False
        dtype: bool
```

```
In [ ]: df2.duplicated().count()
```

```
Out[26]: 50
```

```
In [ ]: df2.duplicated().sum()
```

```
Out[27]: 0
```

```
In [ ]: df2.drop_duplicates(inplace=True)
```

```
In [ ]: df2.duplicated()
```

```
Out[29]: 0      False
1      False
2      False
3      False
4      False
5      False
6      False
7      False
8      False
9      False
10     False
11     False
12     False
13     False
14     False
15     False
16     False
17     False
18     False
19     False
20     False
21     False
22     False
23     False
24     False
25     False
26     False
27     False
28     False
29     False
30     False
31     False
32     False
33     False
34     False
35     False
36     False
37     False
38     False
39     False
40     False
41     False
42     False
43     False
44     False
45     False
46     False
47     False
48     False
49     False
dtype: bool
```

```
In [ ]: df2.dropna()
```


Out[30]:

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84
33	55493.95	103057.49	214634.81	Florida	96778.92
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06

	R&D Spend	Administration	Marketing Spend	State	Profit
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41
49	0.00	116983.80	45173.06	California	14681.40

```
In [ ]: x=df2.iloc[:, :-1].values #features
```

In []: x

```
Out[32]: array([[165349.2, 136897.8, 471784.1, 'New York'],
 [162597.7, 151377.59, 443898.53, 'California'],
 [153441.51, 101145.55, 407934.54, 'Florida'],
 [144372.41, 118671.85, 383199.62, 'New York'],
 [142107.34, 91391.77, 366168.42, 'Florida'],
 [131876.9, 99814.71, 362861.36, 'New York'],
 [134615.46, 147198.87, 127716.82, 'California'],
 [130298.13, 145530.06, 323876.68, 'Florida'],
 [120542.52, 148718.95, 311613.29, 'New York'],
 [123334.88, 108679.17, 304981.62, 'California'],
 [101913.08, 110594.11, 229160.95, 'Florida'],
 [100671.96, 91790.61, 249744.55, 'California'],
 [93863.75, 127320.38, 249839.44, 'Florida'],
 [91992.39, 135495.07, 252664.93, 'California'],
 [119943.24, 156547.42, 256512.92, 'Florida'],
 [114523.61, 122616.84, 261776.23, 'New York'],
 [78013.11, 121597.55, 264346.06, 'California'],
 [94657.16, 145077.58, 282574.31, 'New York'],
 [91749.16, 114175.79, 294919.57, 'Florida'],
 [86419.7, 153514.11, 0.0, 'New York'],
 [76253.86, 113867.3, 298664.47, 'California'],
 [78389.47, 153773.43, 299737.29, 'New York'],
 [73994.56, 122782.75, 303319.26, 'Florida'],
 [67532.53, 105751.03, 304768.73, 'Florida'],
 [77044.01, 99281.34, 140574.81, 'New York'],
 [64664.71, 139553.16, 137962.62, 'California'],
 [75328.87, 144135.98, 134050.07, 'Florida'],
 [72107.6, 127864.55, 353183.81, 'New York'],
 [66051.52, 182645.56, 118148.2, 'Florida'],
 [65605.48, 153032.06, 107138.38, 'New York'],
 [61994.48, 115641.28, 91131.24, 'Florida'],
 [61136.38, 152701.92, 88218.23, 'New York'],
 [63408.86, 129219.61, 46085.25, 'California'],
 [55493.95, 103057.49, 214634.81, 'Florida'],
 [46426.07, 157693.92, 210797.67, 'California'],
 [46014.02, 85047.44, 205517.64, 'New York'],
 [28663.76, 127056.21, 201126.82, 'Florida'],
 [44069.95, 51283.14, 197029.42, 'California'],
 [20229.59, 65947.93, 185265.1, 'New York'],
 [38558.51, 82982.09, 174999.3, 'California'],
 [28754.33, 118546.05, 172795.67, 'California'],
 [27892.92, 84710.77, 164470.71, 'Florida'],
 [23640.93, 96189.63, 148001.11, 'California'],
 [15505.73, 127382.3, 35534.17, 'New York'],
 [22177.74, 154806.14, 28334.72, 'California'],
 [1000.23, 124153.04, 1903.93, 'New York'],
 [1315.46, 115816.21, 297114.46, 'Florida'],
 [0.0, 135426.92, 0.0, 'California'],
 [542.05, 51743.15, 0.0, 'New York'],
 [0.0, 116983.8, 45173.06, 'California']], dtype=object)
```

In []: y=df2.iloc[:, -1].values #labels

```
In [ ]: y
```

```
Out[34]: array([192261.83, 191792.06, 191050.39, 182901.99, 166187.94, 156991.12,
 156122.51, 155752.6 , 152211.77, 149759.96, 146121.95, 144259.4 ,
 141585.52, 134307.35, 132602.65, 129917.04, 126992.93, 125370.37,
 124266.9 , 122776.86, 118474.03, 111313.02, 110352.25, 108733.99,
 108552.04, 107404.34, 105733.54, 105008.31, 103282.38, 101004.64,
 99937.59, 97483.56, 97427.84, 96778.92, 96712.8 , 96479.51,
 90708.19, 89949.14, 81229.06, 81005.76, 78239.91, 77798.83,
 71498.49, 69758.98, 65200.33, 64926.08, 49490.75, 42559.73,
 35673.41, 14681.4 ])
```

```
In [ ]: #now we have to handle the categorical column using column transformer and
```

```
In [ ]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
```

```
In [ ]: from sklearn.preprocessing import OneHotEncoder
ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [3])], remi
x= np.array(ct.fit_transform(x))
```

In []: x

```
Out[39]: array([[0.0, 0.0, 1.0, 165349.2, 136897.8, 471784.1],
 [1.0, 0.0, 0.0, 162597.7, 151377.59, 443898.53],
 [0.0, 1.0, 0.0, 153441.51, 101145.55, 407934.54],
 [0.0, 0.0, 1.0, 144372.41, 118671.85, 383199.62],
 [0.0, 1.0, 0.0, 142107.34, 91391.77, 366168.42],
 [0.0, 0.0, 1.0, 131876.9, 99814.71, 362861.36],
 [1.0, 0.0, 0.0, 134615.46, 147198.87, 127716.82],
 [0.0, 1.0, 0.0, 130298.13, 145530.06, 323876.68],
 [0.0, 0.0, 1.0, 120542.52, 148718.95, 311613.29],
 [1.0, 0.0, 0.0, 123334.88, 108679.17, 304981.62],
 [0.0, 1.0, 0.0, 101913.08, 110594.11, 229160.95],
 [1.0, 0.0, 0.0, 100671.96, 91790.61, 249744.55],
 [0.0, 1.0, 0.0, 93863.75, 127320.38, 249839.44],
 [1.0, 0.0, 0.0, 91992.39, 135495.07, 252664.93],
 [0.0, 1.0, 0.0, 119943.24, 156547.42, 256512.92],
 [0.0, 0.0, 1.0, 114523.61, 122616.84, 261776.23],
 [1.0, 0.0, 0.0, 78013.11, 121597.55, 264346.06],
 [0.0, 0.0, 1.0, 94657.16, 145077.58, 282574.31],
 [0.0, 1.0, 0.0, 91749.16, 114175.79, 294919.57],
 [0.0, 0.0, 1.0, 86419.7, 153514.11, 0.0],
 [1.0, 0.0, 0.0, 76253.86, 113867.3, 298664.47],
 [0.0, 0.0, 1.0, 78389.47, 153773.43, 299737.29],
 [0.0, 1.0, 0.0, 73994.56, 122782.75, 303319.26],
 [0.0, 1.0, 0.0, 67532.53, 105751.03, 304768.73],
 [0.0, 0.0, 1.0, 77044.01, 99281.34, 140574.81],
 [1.0, 0.0, 0.0, 64664.71, 139553.16, 137962.62],
 [0.0, 1.0, 0.0, 75328.87, 144135.98, 134050.07],
 [0.0, 0.0, 1.0, 72107.6, 127864.55, 353183.81],
 [0.0, 1.0, 0.0, 66051.52, 182645.56, 118148.2],
 [0.0, 0.0, 1.0, 65605.48, 153032.06, 107138.38],
 [0.0, 1.0, 0.0, 61994.48, 115641.28, 91131.24],
 [0.0, 0.0, 1.0, 61136.38, 152701.92, 88218.23],
 [1.0, 0.0, 0.0, 63408.86, 129219.61, 46085.25],
 [0.0, 1.0, 0.0, 55493.95, 103057.49, 214634.81],
 [1.0, 0.0, 0.0, 46426.07, 157693.92, 210797.67],
 [0.0, 0.0, 1.0, 46014.02, 85047.44, 205517.64],
 [0.0, 1.0, 0.0, 28663.76, 127056.21, 201126.82],
 [1.0, 0.0, 0.0, 44069.95, 51283.14, 197029.42],
 [0.0, 0.0, 1.0, 20229.59, 65947.93, 185265.1],
 [1.0, 0.0, 0.0, 38558.51, 82982.09, 174999.3],
 [1.0, 0.0, 0.0, 28754.33, 118546.05, 172795.67],
 [0.0, 1.0, 0.0, 27892.92, 84710.77, 164470.71],
 [1.0, 0.0, 0.0, 23640.93, 96189.63, 148001.11],
 [0.0, 0.0, 1.0, 15505.73, 127382.3, 35534.17],
 [1.0, 0.0, 0.0, 22177.74, 154806.14, 28334.72],
 [0.0, 0.0, 1.0, 1000.23, 124153.04, 1903.93],
 [0.0, 1.0, 0.0, 1315.46, 115816.21, 297114.46],
 [1.0, 0.0, 0.0, 0.0, 135426.92, 0.0],
 [0.0, 0.0, 1.0, 542.05, 51743.15, 0.0],
 [1.0, 0.0, 0.0, 0.0, 116983.8, 45173.06]], dtype=object)
```

In []: *#now we can implement the ML model train test split*

In []: **from** sklearn.model_selection **import** train_test_split

```
In [ ]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_st
```

```
In [ ]: x_train
```

```
Out[43]: array([[1.0, 0.0, 0.0, 63408.86, 129219.61, 46085.25],
 [0.0, 1.0, 0.0, 101913.08, 110594.11, 229160.95],
 [0.0, 0.0, 1.0, 78389.47, 153773.43, 299737.29],
 [0.0, 0.0, 1.0, 46014.02, 85047.44, 205517.64],
 [0.0, 0.0, 1.0, 72107.6, 127864.55, 353183.81],
 [0.0, 1.0, 0.0, 91749.16, 114175.79, 294919.57],
 [0.0, 0.0, 1.0, 61136.38, 152701.92, 88218.23],
 [1.0, 0.0, 0.0, 162597.7, 151377.59, 443898.53],
 [0.0, 1.0, 0.0, 93863.75, 127320.38, 249839.44],
 [1.0, 0.0, 0.0, 46426.07, 157693.92, 210797.67],
 [0.0, 0.0, 1.0, 1000.23, 124153.04, 1903.93],
 [0.0, 1.0, 0.0, 75328.87, 144135.98, 134050.07],
 [0.0, 0.0, 1.0, 131876.9, 99814.71, 362861.36],
 [1.0, 0.0, 0.0, 91992.39, 135495.07, 252664.93],
 [0.0, 1.0, 0.0, 73994.56, 122782.75, 303319.26],
 [0.0, 0.0, 1.0, 86419.7, 153514.11, 0.0],
 [0.0, 0.0, 1.0, 94657.16, 145077.58, 282574.31],
 [0.0, 1.0, 0.0, 119943.24, 156547.42, 256512.92],
 [0.0, 1.0, 0.0, 142107.34, 91391.77, 366168.42],
 [0.0, 1.0, 0.0, 27892.92, 84710.77, 164470.71],
 [0.0, 1.0, 0.0, 55493.95, 103057.49, 214634.81],
 [0.0, 0.0, 1.0, 77044.01, 99281.34, 140574.81],
 [1.0, 0.0, 0.0, 100671.96, 91790.61, 249744.55],
 [0.0, 0.0, 1.0, 20229.59, 65947.93, 185265.1],
 [1.0, 0.0, 0.0, 78013.11, 121597.55, 264346.06],
 [0.0, 0.0, 1.0, 542.05, 51743.15, 0.0],
 [0.0, 1.0, 0.0, 1315.46, 115816.21, 297114.46],
 [1.0, 0.0, 0.0, 0.0, 116983.8, 45173.06],
 [0.0, 0.0, 1.0, 120542.52, 148718.95, 311613.29],
 [0.0, 0.0, 1.0, 15505.73, 127382.3, 35534.17],
 [0.0, 0.0, 1.0, 65605.48, 153032.06, 107138.38],
 [1.0, 0.0, 0.0, 64664.71, 139553.16, 137962.62],
 [0.0, 1.0, 0.0, 66051.52, 182645.56, 118148.2],
 [0.0, 0.0, 1.0, 165349.2, 136897.8, 471784.1],
 [0.0, 0.0, 1.0, 114523.61, 122616.84, 261776.23],
 [0.0, 1.0, 0.0, 28663.76, 127056.21, 201126.82],
 [1.0, 0.0, 0.0, 123334.88, 108679.17, 304981.62]], dtype=object)
```

```
In [ ]: y_train
```

```
Out[44]: array([ 97427.84, 146121.95, 111313.02, 96479.51, 105008.31, 124266.9 ,
 97483.56, 191792.06, 141585.52, 96712.8 , 64926.08, 105733.54,
156991.12, 134307.35, 110352.25, 122776.86, 125370.37, 132602.65,
166187.94, 77798.83, 96778.92, 108552.04, 144259.4 , 81229.06,
126992.93, 35673.41, 49490.75, 14681.4 , 152211.77, 69758.98,
101004.64, 107404.34, 103282.38, 192261.83, 129917.04, 90708.19,
149759.96])
```

```
In [ ]: y_train
```

```
Out[45]: array([ 97427.84, 146121.95, 111313.02, 96479.51, 105008.31, 124266.9 ,
                97483.56, 191792.06, 141585.52, 96712.8 , 64926.08, 105733.54,
                156991.12, 134307.35, 110352.25, 122776.86, 125370.37, 132602.65,
                166187.94, 77798.83, 96778.92, 108552.04, 144259.4 , 81229.06,
                126992.93, 35673.41, 49490.75, 14681.4 , 152211.77, 69758.98,
                101004.64, 107404.34, 103282.38, 192261.83, 129917.04, 90708.19,
                149759.96])
```

```
In [ ]: y_test
```

```
Out[46]: array([ 89949.14, 108733.99, 65200.33, 71498.49, 42559.73, 118474.03,
                182901.99, 99937.59, 155752.6 , 156122.51, 81005.76, 191050.39,
                78239.91])
```

```
In [ ]: #now implement the linear regression model to predict the new outcomes
```

```
In [ ]: from sklearn.linear_model import LinearRegression
```

```
In [ ]: reg=LinearRegression()
```

```
In [ ]: reg.fit(x_train,y_train)
```

```
Out[50]: LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [ ]: reg.coef_
```

```
Out[51]: array([ 3.33308401e+02,  1.90646775e+02, -5.23955177e+02,  7.94908963e-01,
                -9.24199896e-03,  2.26310528e-02])
```

```
In [ ]: reg.intercept_
```

```
Out[52]: 49507.17141495356
```

```
In [ ]: reg.score(x_train,y_train)
```

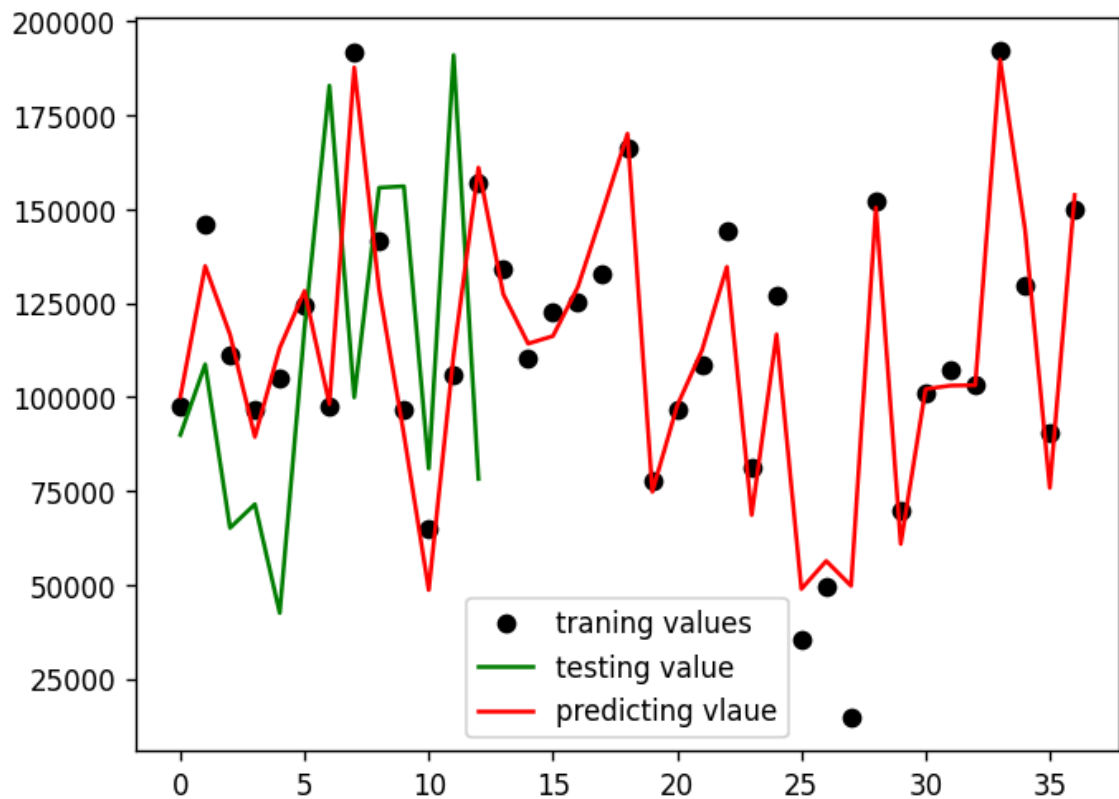
```
Out[53]: 0.9311628200093645
```

```
In [ ]: reg.score(x_test,y_test)
```

```
Out[54]: 0.9874051459541274
```

```
In [ ]: plt.figure(dpi=120)
plt.plot(y_train,"o",color="Black",label="traning values")
plt.plot(y_test,color="Green",label="testing value")
plt.plot(reg.predict(x_train),color="red",label="predicting vlaue")
plt.legend()
```


Out[56]: <matplotlib.legend.Legend at 0x7d597ff94b50>




```
In [ ]: plt.figure(dpi=135)
plt.plot(y, "o")
plt.plot(reg.predict(x)) #best fit line
plt.legend()
```

WARNING:matplotlib.legend:No artists with labels found to put in legend.
Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

Out[57]: <matplotlib.legend.Legend at 0x7d598000e080>



```
In [ ]: df3=pd.DataFrame() #need to check
df3["ActualValue"]=y_test
df3["PredictedValue"]=reg.predict(x_test)
print(df3)
```

	ActualValue	PredictedValue
0	89949.14	88857.102563
1	108733.99	109299.917912
2	65200.33	66680.290487
3	71498.49	71093.303448
4	42559.73	48588.864362
5	118474.03	116162.086537
6	182901.99	171321.584718
7	99937.59	99971.425331
8	155752.60	159257.651210
9	156122.51	158377.469861
10	81005.76	83684.483035
11	191050.39	179967.050876
12	78239.91	75512.499927

In []:

